



# **Bahir Dar Univerisity**

**Bahir Dar Institute Of Technology  
Faculity Of Computing  
Deparrtment Of Software Engineering  
Operating System And System Programming  
Individual Assignment 2**

**Title : Splice system call**

`ssize_t splice(int fd_in, loff_t *off_in, int fd_out, loff_t *off_out, size_t len, unsigned int flags)`

**Name:**

**ID**

**Fentahun Mengie**

**1306919**

**Submitted to:InstructorWendimu Baye  
Submission date: 17-11-2014 E.C**

# Introduction

`ssize_t splice(int fd_in, loff_t *off_in, int fd_out, loff_t *off_out, size_t len, unsigned int flags)`

In this document we will see what is splice system call?, why we use splice system call?, how we use system call and its parameters and its flags. Finally we will see code implementation of the splice system call.

**splice()** is a Linux-specific system call that moves data between two file descriptors without copying between kernel address space and user address space. It moves data between a file descriptor and a pipe without a round trip to user space.

It transfers up to len bytes of data from the file descriptor fd\_in to the file descriptor fd\_out, where one of the file descriptors must refer to a pipe.

With splice(), data can be transferred from one file descriptor to another without requiring any copies from user space into kernel space, which is typically necessary **to maintain system security** and to **provide processes with a straightforward interface** for reading and writing to files. splice() works by using the pipe buffer. A pipe buffer is an in-kernel memory buffer that is opaque to the user space process. A user process can splice the contents of a source file into this pipe buffer, then splice the pipe buffer into the destination file, all without moving any data through userspace.

## Parameters and Flags

- ❖ **fd\_in** and **fd\_out** represent the input and output file descriptors, respectively, and one of these two file descriptors must point to a pipeline device, which is a less friendly restriction.
- ❖ **off\_in** and **off\_out** are pointers to the offsets of **fd\_in** and **fd\_out** respectively, indicating where the kernel reads and writes data from.

The following semantics apply for fd\_in and off\_in:

If fd\_in refers to a pipe, then off\_in must be NULL.

If fd\_in does not refer to a pipe and off\_in is NULL, then bytes are read from fd\_in starting from the file offset, and the file offset is adjusted appropriately.

If fd\_in does not refer to a pipe and off\_in is not NULL, then off\_in must point to a buffer which specifies the starting offset from which bytes will be read from fd\_in; in this case, the file offset of fd\_in is not changed.

Analogous statements apply for fd\_out and off\_out.

- ❖ **len** indicates the number of bytes the call wishes to transfer

- ❖ **Flag** is the system call's flag option bitmask, which sets the behavior of the system call, and is a combination of 0 or more of the following values via the 'or' operation.

**SPLICE\_F\_MOVE**

**SPLICE\_F\_NONBLOCK**

**SPLICE\_F\_MORE**

**SPLICE\_F\_GIFT**

- **SPLICE\_F\_MOVE**: instructs splice() to try to just move memory pages instead of copying them; setting this value does not necessarily mean that memory pages will not be copied; whether they are copied or moved depends on whether the kernel can move memory pages from the pipeline, or whether the memory pages in the pipeline are intact. pages may be copied if the kernel can not move the pages from the pipe, or if the pipe buffers don't refer to full pages. The initial implementation of this flag had a lot of bugs, so it has been in place since Linux version 2.6.21, but it has been retained because it may be reimplemented in a future version.
- **SPLICE\_F\_NONBLOCK**: instructs splice() not to block I/O, i.e. makes the splice() call a non-blocking call that can be used to implement asynchronous data transfers, but note that it is also best to pre-mark the two file descriptors for data transfers as non-blocking I/O with O\_NONBLOCK, Otherwise, the splice() call may still be blocked.
- **SPLICE\_F\_MORE**: informs the kernel that more data will be transferred with the next splice() system call, this flag is useful for scenarios where the output/fd\_out/ side is a socket.
- **SPLICE\_F\_GIFT**: Unused for splice()

## RETURN VALUE

Upon successful completion, splice() returns the number of bytes spliced to or from the pipe. A return value of 0 means end of input.

If fd\_in refers to a pipe, then this means that there was no data to transfer, and it would not make sense to block because there are no writers connected to the write end of the pipe.

On error, splice() returns -1 and errno is set to indicate the error.

## ERRORS

**EAGAIN SPLICE\_F\_NONBLOCK** was specified in flags or one of the file descriptors had been marked as nonblocking (O\_NONBLOCK), and the operation would block.

**EBADF** One or both file descriptors are not valid, or do not have proper read-write mode.

**EINVAL**

- ✓ The target filesystem doesn't support splicing.
- ✓ The target file is opened in append mode
- ✓ Neither of the file descriptors refers to a pipe.
- ✓ An offset was given for nonseekable device (e.g., a pipe).
- ✓ fd\_in and fd\_out refer to the same pipe.

**ENOMEM** Out of memory.

**ESPIPE** Either off\_in or off\_out was not NULL, but the corresponding file descriptor refers to a pipe.

## code implementation

```
#define _GNU_SOURCE
#include <errno.h>
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/stat.h>

int main(int argc, char **argv) {
    if(argc < 2) {
        fputs("it is copied succesfully\n", stderr);
        return 1;
    }

    FILE *handle = fopen(argv[1], "r");
    if(!handle) {
        perror("Error opening file:");
        goto error;
    }

    int handle_fd = fileno(handle);

    int filedes[2];
    if(pipe(filedes)) {
        perror("Error creating pipe:");
        goto error;
    }

    while(1) {
        ssize_t rc = splice(handle_fd, NULL, filedes[1], NULL, BUFSIZ, 0);

        if(rc == -1) {
            // Error occurred
            perror("Error copying data:");
            goto error;
        } else if(rc == 0) {
            break;
        }
    }
}
```

```
    }

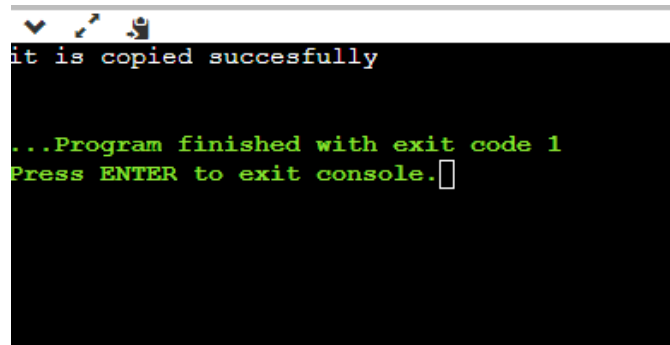
    splice(filedes[0], NULL, STDOUT_FILENO, NULL, BUFSIZ, 0);
}

return 0;

error:
    if(fclose(handle)) {
        perror("Error closing file:");
    }

    return 1;
}
```

Here is the output. As we can see from below it is copied  
I hav run it in online c compiler

A screenshot of a terminal window from an online C compiler. The terminal has a black background with white and green text. At the top, there are three small icons: a checkmark, a cursor, and a document. The text in the terminal reads: "it is copied succesfully" in white, followed by "...Program finished with exit code 1" in green, and "Press ENTER to exit console." in green with a white cursor at the end.