BIT

Bahir Dar University institute of Technology

Operating System and System Programming Individual second assignment

On System Call

Prepared by: Tadele Workie Mihretu.

ID: 1308433

```
int getdents(unsigned int fd, struct linux_dirent *dirp,
unsigned int count);
```

Submitted to Instructor Wondimu Baye.

Submission Date: JULY 17, 2014 E.C

CONTENTS

Introduction **3**

What / why / how, getdents system call? **3**

Briefly describe about the list of parameters. **4**

LIst of flags and their implementation **5**

Source code implementation using O_RDONLY flag 6

The output of the above source code using the O_RDONLY flag 8

The source code using O_DIRECTORY flag 8

The output of the above source code using O_DIRECTORY 10

Source code implementation by using both flags at the same time. 10

The output of the above source code: 12

Reference: **12**

INTRODUCTION

System call is a mechanism that provides the interface between a process and the operating system. It is a programmatic method in which a computer program requests a service from the kernel of the OS.System call offers the services of the operating system to the user programs via API (Application Programming Interface). System calls are the only entry points for the kernel system. A system call **connects to the operating system's kernel, which executes in kernel space**. When an application creates a system call, it must first obtain permission from the kernel. It achieves this using an interrupt request, which pauses the current process and transfers control to the kernel.

Syntax of getdents () system call.

```
#include <sys/syscall.h>      /* Definition of SYS_* constants *
#include <unistd.h>

int getdents(unsigned int fd, struct linux_dirent *dirp,
unsigned int count);
```

WHAT / WHY / HOW, GETDENTS SYSTEM CALL?

The system call **getdents** () reads several *linux_dirent* structures from the directory referred to by the open file descriptor *fd* into the buffer pointed to by *dirp*. The argument *count* specifies the size of that buffer.

The getdents() function attempts to read nbyte bytes from the directory associated with the file descriptor fildes and to format them as file system independent directory entries in the buffer pointed to by buffer. Since the file system independent directory entries are of variable lengths, in most cases the actual number of bytes returned will be less than nbyte. The file system independent directory entry is specified by the dirent structure.

struct linux dirent will be returned by getdents. It will do this for any underlying file system type. Because the "on disk" format may be completely different, known only to the given file system driver, a simple user space read call may fail. To put it another way, getdents can convert from the native format to fill the linux dirent.

Directory entries in UNIX can refer to files, but also to directories, named pipes and devices. The character and block device entries are the interface to the different drivers in the kernel. They contain a number to identify the driver, and another number to identify different devices handled by the same driver.

Newer systems also support file system sockets, but the whole socket concept is not part of the original UNIX design. Sockets are not files, and "everything is a file" refers to the original UNIX design. But once sockets are set up, they support some operations that also work on files.

RETURN VALUE

On success, the number of bytes read is returned. On end of direc-tory, 0 is returned.
On error, -1 is returned, and errno is set ap-propriately.

ERRORS

EBADF Invalid file descriptor fd.

EFAULT Argument points outside the calling process's address space.

EINVAL Result buffer is too small.

ENOENT No such directory.

ENOTDIR File descriptor does not refer to a directory.

BRIEFLY DESCRIBE ABOUT THE LIST OF PARAMETERS.

There getdents() has three parameters with during this implementation.

- ***fd***
- **count**
- **struct linux_dirent *dirp**

fd :- file descriptor it is used to access the file.

Count: - it is used to describe the size of buffer

The linux_dirent structure : is declared as follows:

```
struct linux_dirent {

    unsigned long  d_ino;
    unsigned long  d_off;
    unsigned short d_reclen;
    char           d_name[];
    char           pad;
    char           d_type;
};
```

d_ino *is an inode number.*
d_off *is the distance from the start of the directory to the start of the next linux_dirent.*
d_reclen *is the size of this entire linux_dirent.*
d_name *is a null-terminated filename.*
d_type *is a byte at the end of the structure that indicates the file type.*
It contains one of the following values (defined in <di-rent.h>):

DT_BLK This is a block device.

DT_CHR This is a character device.

DT_DIR This is a directory.

DT_FIFO *This is a named pipe (FIFO).*

DT_LNK *This is a symbolic link.*

DT_REG *This is a regular file.*

DT_SOCK *This is a UNIX domain socket.*

DT_UNKNOWN *The file type is unknown.*

There is a two flags also in getdents system call

- **O_RDONLY**

We can use the "O_RDONLY" flag of the C programming language only if we have included the "sys/types.h", "sys/stat.h", and "fcntl.h" header files in our C script. In this simple C program, we have defined an integer type variable "fd" that refers to the file descriptor of the file that we want to open as read-only. Then, we have used the "open()" function of the C programming language and have passed to it the path of the desired file followed by the "O_RDONLY" flag indicating that we want to open the file as read-only. Finally, we have printed a confirmation message on the terminal using the "printf" statement.

- **O_DIRECTORY**

If pathname is not a directory, cause the open to fail. This flag was added in kernel version 2.1.126, to avoid denial-of-service problems if **open()** is called on a FIFO or tape device.

3. List the flags, their purpose with code implementation (give example source code with output)

LIST OF FLAGS AND THEIR IMPLEMENTATION

Source code implementation using O_RDONLY flag

```
#define _GNU_SOURCE
#include <dirent.h>
#include <fcntl.h>
#include <stdint.h>
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
#include <sys/stat.h>
#include <sys/syscall.h>
#define handle_error(msg)
               do { perror(msg); exit(EXIT_FAILURE); } while (0)

       struct linux_dirent {
           unsigned long  d_ino;
           off_t          d_off;
           unsigned short d_reclen;
           char           d_name[];
       };

       #define BUF_SIZE 1024

       int
       main(int argc, char *argv[])
       {
           int fd;
           long nread;
           char buf[BUF_SIZE];
           struct linux_dirent *d;
           char d_type;

           fd = open(argc > 1 ? argv[1] : ".", O_RDONLY);
           if (fd == -1)
               handle_error("open");

           for (;;) {
               nread = syscall(SYS_getdents, fd, buf, BUF_SIZE);
               if (nread == -1)
                   handle_error("getdents");

               if (nread == 0)
                   break;

               printf("--------------- nread=%d ---------------\n", nread);
               printf("inode#    file type  d_reclen  d_off   \t\t d_name\n");
```

```
            for (long bpos = 0; bpos < nread;) {
                d = (struct linux_dirent *) (buf + bpos);
                printf("%8ld  ",  d->d_ino);
                d_type = *(buf + bpos + d->d_reclen - 1);
                printf("%-10s ", (d_type == DT_REG) ?  "regular" :
                                 (d_type == DT_DIR) ?  "directory" :
                                 (d_type == DT_FIFO) ? "FIFO" :
                                 (d_type == DT_SOCK) ? "socket" :
                                 (d_type == DT_LNK) ?  "symlink" :
                                 (d_type == DT_BLK) ?  "block dev" :
                                 (d_type == DT_CHR) ?  "char dev" : "???");
                printf("%4d %10jd  %s\n", d->d_reclen,
                        (intmax_t) d->d_off, d->d_name);
                bpos += d->d_reclen;
            }
        }
        exit(EXIT_SUCCESS);
}
```

The output of the above source code using the O_RDONLY flag

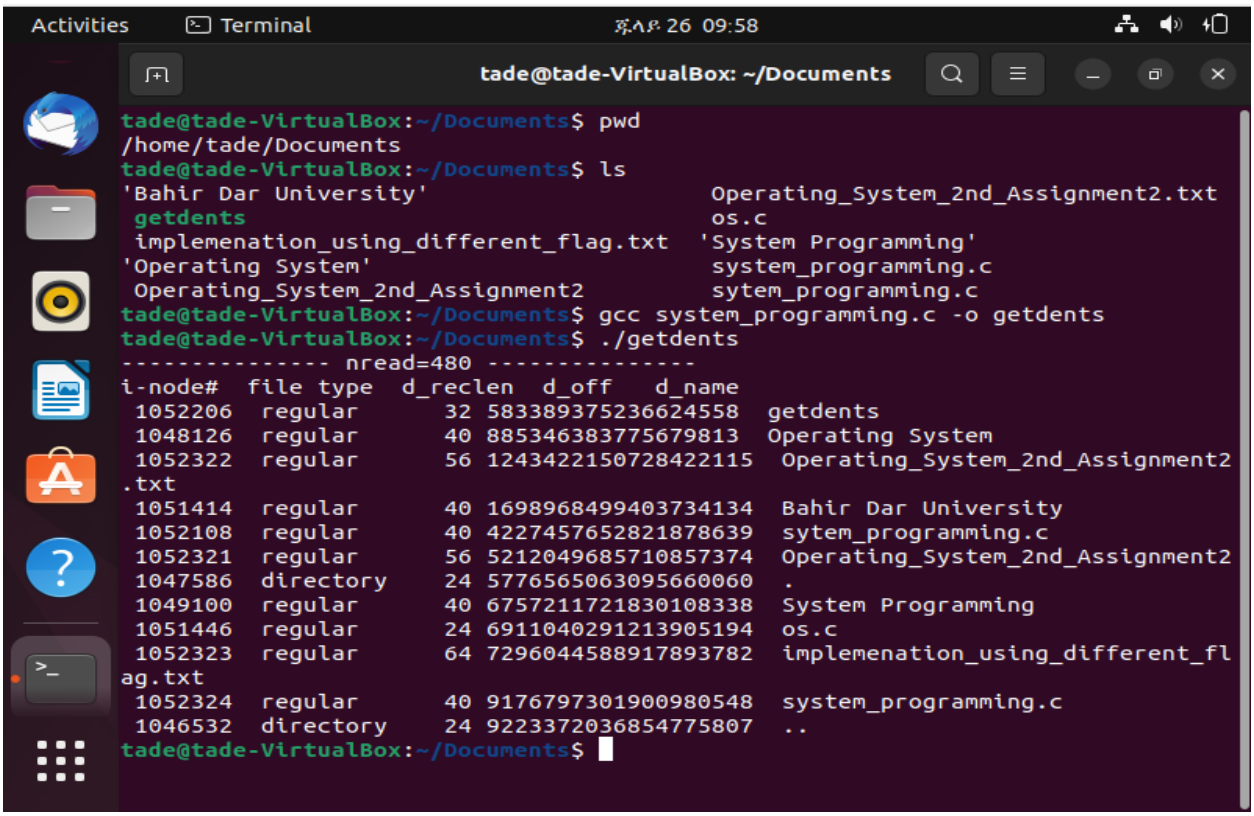

The source code using O_DIRECTORY flag

```c
#define _GNU_SOURCE
#include <dirent.h>
#include <fcntl.h>
#include <stdint.h>
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
#include <sys/stat.h>
#include <sys/syscall.h>
#define handle_error(msg)
               do { perror(msg); exit(EXIT_FAILURE); } while (0)

       struct linux_dirent {
           unsigned long  d_ino;
           off_t          d_off;
           unsigned short d_reclen;
           char           d_name[];
```

```
    };

    #define BUF_SIZE 1024

    int
    main(int argc, char *argv[])
    {
        int fd;
        long nread;
        char buf[BUF_SIZE];
        struct linux_dirent *d;
        char d_type;

        fd = open(argc > 1 ? argv[1] : ".", O_RDONLY);
        if (fd == -1)
            handle_error("open");

        for (;;) {
            nread = syscall(SYS_getdents, fd, buf, BUF_SIZE);
            if (nread == -1)
                handle_error("getdents");

            if (nread == 0)
                break;
            printf("--------------- nread=%d ---------------\n", nread);
            printf("inode#    file type  d_reclen  d_off   \t\t d_name\n");
            for (long bpos = 0; bpos < nread;) {
                d = (struct linux_dirent *) (buf + bpos);
                printf("%8ld  ",  d->d_ino);
                d_type = *(buf + bpos + d->d_reclen - 1);
                printf("%-10s ", (d_type == DT_REG) ?  "regular" :
                                 (d_type == DT_DIR) ?  "directory" :
                                 (d_type == DT_FIFO) ? "FIFO" :
                                 (d_type == DT_SOCK) ? "socket" :
                                 (d_type == DT_LNK) ?  "symlink" :
                                 (d_type == DT_BLK) ?  "block dev" :
                                 (d_type == DT_CHR) ?  "char dev" : "???");
                printf("%4d %10jd  %s\n", d->d_reclen,
                        (intmax_t) d->d_off, d->d_name);
                bpos += d->d_reclen;
            }
        }
        exit(EXIT_SUCCESS);
}
```

The output of the above source code using O_DIRECTORY

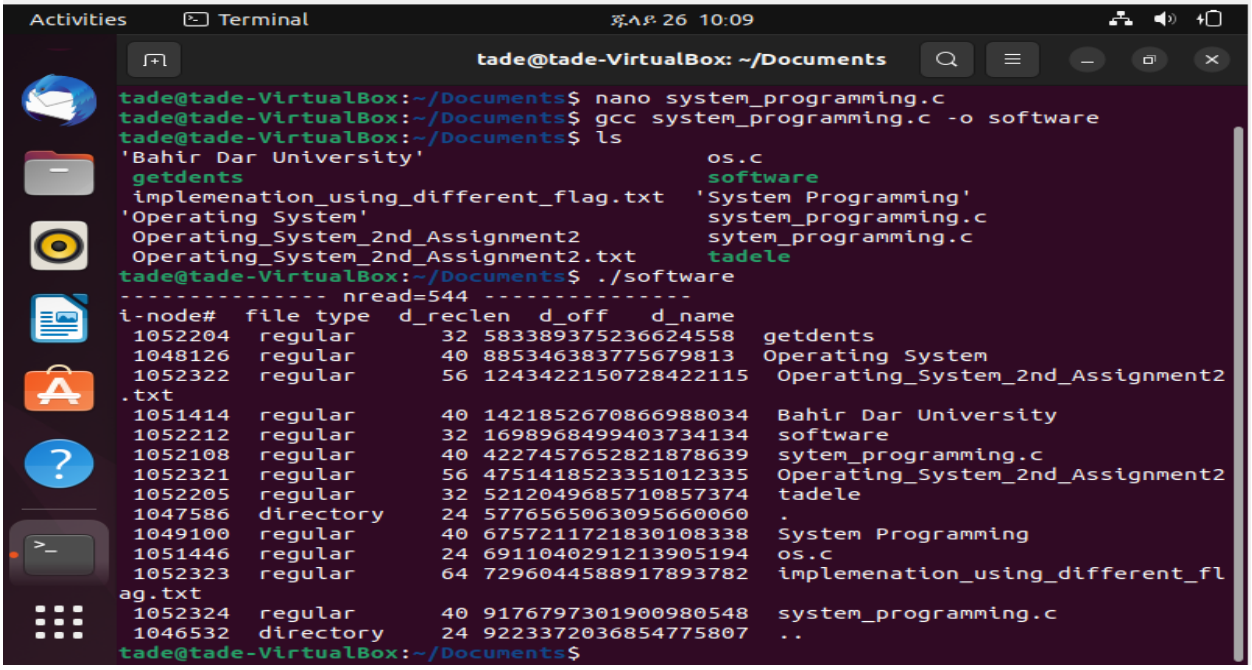

Source code implementation by using both flags at the same time.

```c
#define _GNU_SOURCE
#include <dirent.h>
#include <fcntl.h>
#include <stdint.h>
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
#include <sys/stat.h>
#include <sys/syscall.h>
#define handle_error(msg)
               do { perror(msg); exit(EXIT_FAILURE); } while (0)

       struct linux_dirent {
           unsigned long  d_ino;
           off_t          d_off;
           unsigned short d_reclen;
           char           d_name[];
       };
```

```c
#define BUF_SIZE 1024

int
main(int argc, char *argv[])
{
    int fd;
    long nread;
    char buf[BUF_SIZE];
    struct linux_dirent *d;
    char d_type;

    fd = open(argc > 1 ? argv[1] : ".", O_RDONLY | O_DIRECTORY);
    if (fd == -1)
        handle_error("open");

    for (;;) {
        nread = syscall(SYS_getdents, fd, buf, BUF_SIZE);
        if (nread == -1)
            handle_error("getdents");

        if (nread == 0)
            break;

        printf("--------------- nread=%d ---------------\n", nread);
        printf("inode#    file type  d_reclen  d_off   \t\t d_name\n");
        for (long bpos = 0; bpos < nread;) {
            d = (struct linux_dirent *) (buf + bpos);
            printf("%8ld  ",  d->d_ino);
            d_type = *(buf + bpos + d->d_reclen - 1);
            printf("%-10s ", (d_type == DT_REG) ?  "regular" :
                             (d_type == DT_DIR) ?  "directory" :
                             (d_type == DT_FIFO) ? "FIFO" :
                             (d_type == DT_SOCK) ? "socket" :
                             (d_type == DT_LNK) ?  "symlink" :
                             (d_type == DT_BLK) ?  "block dev" :
                             (d_type == DT_CHR) ?  "char dev" : "???");
            printf("%4d %10jd  %s\n", d->d_reclen,
                    (intmax_t) d->d_off, d->d_name);
            bpos += d->d_reclen;
        }
    }
    exit(EXIT_SUCCESS);
}
```

The output of the above source code:

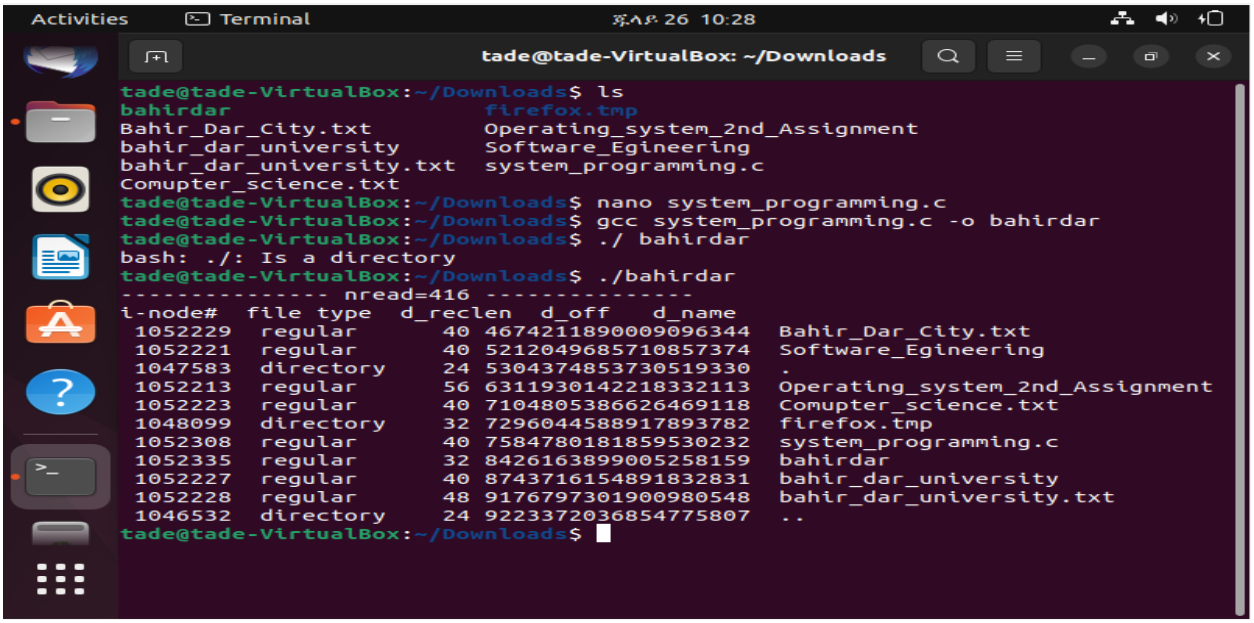

REFERENCE:

The Linux and Unix System programming Handbook
https://man7.org/linux/man-pages/man2/getdents.2.html.
https://linuxhint.com/c-language-o_donly-o_wrongly-and-o_rdwr-flags/